

1 Method

The proposed model minimizes a composite objective that combines a pixel term with a structural term, weighted so that neither dominates the optimisation during the early epochs of training.

$$\mathcal{L} = \sum_{k=1}^3 \lambda_k \left(\frac{\|x_k - y_k\|_2^2}{\sqrt{N_k}} \right) + \epsilon. \quad (1)$$

Here λ_k weights each scale and N_k is the number of pixels at scale k . The architecture is summarised in Fig. ?? . The re-




Figure 1: Overview of the network.

maining sections describe the training schedule and the evaluation protocol in detail, including the data split, the augmentation policy and the metrics used to compare against the baselines.